

VICTORIA UNIVERSITY OF WELLINGTON
Te Whare Wananga o te Upoko o te Ika a Maui



Cassandra CQL Queries

Lecturer : Dr. Pavle Mogin

SWEN 432
*Advanced Database Design and
Implementation*

Plan for CQL Queries

- The Syntax of the `SELECT` Statement
 - The Syntax of the `select_expression`
 - The Syntax of the `relation` in the `WHERE` Clause
- Simple `SELECT` expressions
- Filtering Data using `WHERE` Clause
- Restrictions on Conditions in the `WHERE` Clause
- Using Indexes
- Filtering Collections
- Querying Tables with Columns of the `counter` Type
- Keyspace Design Heuristics
 - **Readings:** Have a look at *Useful Links at the Course Home Page*

The Syntax of CQL `SELECT` Statement

```
SELECT select_expression
FROM keyspace_name.table_name
WHERE relation AND relation ...
ORDER BY (clustering_column (ASC | DESC) ...)
LIMIT n
ALLOW FILTERING
```

The Syntax of `select_expression`

`select_expression` is:

```
selection_list | DISTINCT selection_list |  
(COUNT (* | 1) )
```

`selection_list` is one of:

- A list of partition keys (used with `DISTINCT`)
- `selector AS alias, selector AS alias, ...| *`
 - `alias` is an alias for a column name
- A selector is:

```
column_name | ( WRITETIME (column_name) ) |  
( TTL (column_name) ) | (function (selector,  
selector, ... ) )
```

- `function` is a `timeuuid` function, a `token` function, or a `blob` conversion function

The Syntax of relation

- A relation is:

```
column_name op term |  
column_name IN ( term, ( term ... ) ) |  
TOKEN (column_name, ...) op ( term )
```

- An op is:

```
= | < | > | <= | > | = | CONTAINS | CONTAINS KEY
```

- A term is:

- A constant: string, number, uuid, boolean, hex
- A bind marker (?)
- A function
- A collection: set: {literal, ...}, list: [literal, ...],
map: {literal: literal, ...}

Using CQL *SELECT* Statement

- The CQL *SELECT* statement works very similar to the SQL *SELECT* statement with a few exceptions
 - The main exception is that it retrieves one or more records from a single table only
 - The others will be discussed through examples

- Simple examples

```
use blogs;
```

```
SELECT * FROM users;
```

```
SELECT COUNT(*) FROM blog_entries;
```

```
SELECT DISTINCT user_name, body  
FROM blog_entries;
```

- The last *SELECT* returns only one row for each distinct *user_name* that is the partition key

More Simple *SELECT* Examples

```
ALTER TABLE users ADD created_at timeuuid;
```

```
INSERT INTO users (user_name, name,  
created_at) VALUES ('jbond', 'James',  
now());
```

```
// now() returns a unique timeuuid in ms
```

```
SELECT user_name, name, dateOf(created_at)  
AS creation_date FROM users;
```

```
//dateOf() returns the timestamp part of an  
//timeuuid
```

Filtering Data Using *WHERE*

- In the `WHERE` clause, columns should be referred using the actual names, not aliases
- Each column in the `WHERE` clause has:
 - Either to belong to the partition key or
 - To be indexed using `CREATE INDEX`
- A `WHERE` clause relation has to be build by putting:
 - The name of the column to the left of an operator and
 - The column value to the right of the operator

```
SELECT * FROM users  
WHERE user_name = 'jbond';
```

- Note, `user_name` is the partitioning column

Restrictions on Conditions

(1)

- Cassandra does not support non-equal conditional operations on the partition key
 - The token function should be used for range queries on the partition key

```
SELECT * FROM users WHERE  
TOKEN(user_name) >= TOKEN('asmith') AND  
TOKEN(user_name) =< TOKEN('jbond');
```

Restriction on The Use of Conditions (2)

- The `IN` condition is allowed on the last column of the partition key only if the query contains equality conditions on all preceding columns of the key
- Assume, the `blog_entries` table has been defined in the following way:

```
CREATE TABLE blog_entries (user_id text,  
date int, body text, no int,  
PRIMARY KEY ((user_name, date), no));
```

```
SELECT * FROM blog_entries WHERE  
user_name = 'jbond' AND date IN  
(20150810, 20150811, 20150812);
```

The ALLOW FILTERING Clause

- When one attempts a potentially expensive query, such as searching a range of rows, Cassandra prompts the following message:

```
Bad Request: Cannot execute this query as  
it might involve data filtering and thus  
may have unpredictable performance. If you  
want to execute this query despite the  
performance unpredictability, use ALLOW  
FILTERING.
```

- In such cases, imposing a limit using the `LIMIT n` clause is also recommended to reduce memory used

Restriction on The Use of Conditions (3)

- Cassandra supports greater-than and less-than comparisons on a clustering column, but for a given partition key, the conditions on the clustering column are restricted to the filters that allow Cassandra to select a contiguous ordering of rows
- Assume `blog_entries` primary key is `(user_name, date)`

```
SELECT * FROM blog_entries WHERE  
user_name = 'jbond' AND  
date >= 20150810 AND date < 20160410  
ALLOW FILTERING;
```

Secondary Indexes

- Building and using a secondary index gives best results if the filtering column has many duplicate values
- Cassandra would refuse to execute a query asking for users from a certain city unless `ALLOW FILTERING` has been defined

```
CREATE INDEX city_index on users(city);
```

```
SELECT user_name, name FROM users  
WHERE city = 'London' LIMIT 10;
```

- If a secondary index on the filtering field has been defined, the `ALLOW FILTERING` clause is not needed
- If a secondary index on the filtering field has not been defined, the `ALLOW FILTERING` clause is needed

Filtering Collections

- Cassandra retrieves the collection in its entirety
 - Assume `subscribers` column of the `subscribed_to` table is of the type `set`

```
SELECT user_name, subscribers
FROM subscribed_to;
```

- A table containing a collection column can be indexed on the collection column
- The `CONTAINS` condition in the `WHERE` clause can be used to filter the data for a particular value in the collection

```
CREATE INDEX ON subscribed_to
(subscribers);
```

```
SELECT user_name FROM subscribed_to
WHERE subscribers CONTAINS 'canslow';
```

Filtering a map Collection Column

- Assume now, subscribers column of the subscribed_to table is of the type map
 - This is an index on map values:

```
CREATE INDEX my_map_values_index ON  
subscribed_to (subscribers);
```

```
SELECT user_name FROM subscribed_to WHERE  
subscribers CONTAINS 'London';
```

- Indexes on keys and values of a map can't co-exist

```
DROP INDEX my_map_values_index;
```

```
CREATE INDEX my_map_keys_index ON  
subscribed_to (KEYS(subscribers));
```

```
SELECT user_name FROM subscribed_to  
WHERE subscribers CONTAINS KEY 'canslow';
```

Tables With counter Data Type

- Extending `blogs` key space by a table for counting daily and total number of visits to users' blogs

```
CREATE TABLE visits (user_name text, date  
int, total_count counter static,  
daily_count counter, PRIMARY KEY(  
user_name, date);
```

```
UPDATE visits SET dayly_count =  
daily_count + 1, total_count = total_count  
+ 1 WHERE user_name = 'jbond' AND date =  
20160317;
```

```
SELECT DISTINCT total_count FROM visits  
WHERE user_name = 'jbond';
```


COUNT (*) versus counter Tables

- Assume the `blog_entries` table has the primary key `((user_name, date), no)`
- To find the total number of blogs of the user `jbond`, the following query has to search through several partitions

```
SELECT COUNT(*) FROM blog_entries WHERE  
user_name = 'jbond';
```

- A query to a table having a counter column containing the number of blogs for each user, would require accessing just one partition

Queries to *blog* Keyspace

- Each table in a Cassandra keyspace is aimed for a few specific queries
- In the `blog` keyspace:
 - The table `users` is aimed for answering queries about bloggers data,
 - The table `blog_entries` is aimed for retrieving blogs of a user,
 - The table `subscribes_to` is aimed for answering the question: who are subscribers to a user's blog
 - The table `visits` is aimed for retrieving counts of blogs made by a user
- Each table represents a materialized view for answering at least one query
 - Loosely speaking the answers for all anticipated queries are already physically stored

Database Design Principle Guidelines

- Each table of a keyspace has a self-contained set of columns intended to support one (or possibly more) specific queries
 - Since joins are not supported, for each query there has to exist at least one table containing all data needed to satisfy the query
 - Additionally, data representing the query condition need to belong either to partition key or to be indexed, otherwise `ALLOW FILTERING` clause has to be used
- The statements above are the main guidelines for designing tables implying that the starting point in designing tables represents the identification of user queries

Formalized Database Design Guidelines

- Let q be a query having:
 - The set of column names $C(s)$ in its `SELECT` clause and
 - The set of column names $C(w)$ in its `WHERE` clause
- Let t be a table having:
 - The set of column names $C(t)$,
 - The set of partition key column names $C(k)$, and
 - The set of indexed column names $C(i)$
- If

$$(C(t) \supseteq (C(s) \setminus \{\text{COUNT}(\ast)\})) \wedge ((C(k) \cup C(i)) \supseteq C(w))$$
 then the query q can be efficiently satisfied by table t
- If

$$(\exists a \in \alpha(t) \cap \alpha(w)) (a \notin (\alpha(k) \cup \alpha(i)))$$
 the clause `ALLOW FILTERING` has to be used

Other Design Considerations

- The primary key values have to be unique
- The primary key can be a superkey
 - The primary key can contain logically redundant columns
 - To avoid creating secondary indexes
 - To use the same table for answering more than one query efficiently

- Consider queries q_1 and q_2 , having

$$(C(w_1) \subseteq C(w_2))$$

- Both queries can be efficiently satisfied by a table t having

$$((C(k) \cup C(i)) \supseteq C(w_2)) \wedge (C(t) \supseteq (C(s_1) \cup C(s_2)))$$

Summary

- **Syntax of the CQL SELECT statement:**

```
SELECT select_expression
FROM keyspace_name.table_name
WHERE relation AND relation ...
ORDER BY (clustering_column (ASC | DESC) ...)
LIMIT n
ALLOW FILTERING
```

- **The `relation` clause should contain:**
 - Either primary key and indexed columns, or
 - The `ALLOW FILTERING` clause has to be used
- Since joins are not supported, tables are designed to contain all data needed to answer a set of queries