

VICTORIA UNIVERSITY OF WELLINGTON  
*Te Whare Wananga o te Upoko o te Ika a Maui*



# ***Cassandra Storage Engine***

***Lecturer : Dr. Pavle Mogin***

SWEN 432  
*Advanced Database Design and  
Implementation*

# ***Cassandra The Fortune Teller***

---



# ***Plan for Cassandra Storage Engine***

---

- Table Primary Key and Partitioning
- Storage Engine Rows
- Log Structured Merge Trees (LSM Trees)
- Memtable and SSTables
- Write Paths for Insert and Update
- About Reads
- About Deletes
- Compaction:
  - Size Tiered Compaction
  - Leveled Compaction
- ***Readings:*** Have a look at *Useful Links at the Course Home Page*

# ***Table Primary Key and Partitioning (1)***

---

- The table primary key is defined within the table declaration
- Each table row has a unique primary key value
- A primary key can be:
  - Simple containing just one column name, where the column name is the partition key, or
  - Compound containing more than one column name, where the first column name is the partition key and the remaining column names are the clustering key, or
  - Composite having a multicolumn partition key enclosed in a parenthesis, and a clustering key
- Example:
  - The table `users` has a simple primary key `user_name`,
  - The table `blogs_entry` has a composite primary key `((user_name, date), no)`

# ***Table Primary Key and Partitioning (2)***

---

- All table rows having the same partition key value make a CQL partition
  - A single primary key makes single row CQL partitions,
  - A compound primary key and a composite primary key have both the partition and clustering keys and produce multi row CQL partitions
  - CQL rows within a partition are sorted according to the clustering key values
- Table rows are assigned to nodes in the consistent hashing ring by
  - A cluster configured partitioner and
  - The replica placement strategy defined for each keyspace
- The partitioner hashes partition keys into tokens
  - Tokens are points on the consistent hashing ring
  - This way are CQL partitions assigned to cluster nodes

# ***Indexes***

---

- The primary index of a table is a unique index on its row key
  - Cassandra maintains the primary index automatically
- Each node maintains this index for data it manages

# Storage Engine Row

---

- A storage engine row is a sequence of table rows having the same partition key value stored on disk
  - For tables having a primary key with no clustering column, a storage engine row contains a single table row
  - For tables having a primary key with clustering columns, a storage engine row contains at least one table row
- As a partition size grows, the storage engine row grows
- As the storage engine row grows, read latency for some queries will increase
- Example:
  - Assume `blog_entries` table key is `(user_name, no)`. As new blogs are added into `blog_entries` table, storage engine rows for particular users may grow very big

# Bucketing

---

- A technique to limit the size of large engine rows is to introduce a sensible time bucket
- Example:
  - In the case of the `blog_entries` table a sensible time bucket may be a `year_month` column that extends the primary key in the following way  
(`(user_name, year_month), no`)
- But, there is no way to change a primary key in Cassandra, as it defines how data is stored physically
- The only work around is:
  - To create a new table with the new primary key,
  - Copy data from the old one, and then
  - Drop the old table
- Bucketing has to be defined at the time of the database design



# Reverse Queries

---

- Often, queries ask for the last entry of a time series
- Example:
  - Retrieve the James's last blog
- One way to satisfy the query is to sort the table for each query:

```
SELECT * FROM blog_entries  
WHERE user_name = 'jbond'  
ORDER BY no desc  
LIMIT 1;
```

- Read performance wise, a more efficient way is to keep CQL partitions sorted in the descending order

# Redesigning the *blog\_entries* Table

---

- Table:


```
CREATE TABLE blog_entries (  
  user_name text,  
  year_month int  
  body text,  
  no int  
  PRIMARY KEY ((user_name, year_month), no)  
  ) with clustering order by (no desc);
```

- Query:

```
SELECT * FROM blog_entries WHERE  
user_name = 'jbond' and year_month =  
201603 LIMIT 1;
```


# Log Structured Merge Trees

- LSM trees are an approach to use memory and disk storage to satisfy write and read requests in an efficient and safe way
  - In the memory, there is a memtable containing chunks of recently committed data,
  - On disk, there is a commit-log file and a number of SSTables containing data flushed from the memtable




Table's data reside in the Memtable and SSTable 1, SSTable 2, and SSTable 3

SSTables are *immutable*




# LSM Trees Write and Read Paths

## Write Path



LSM trees are optimized for writes since writes go only to Commit Log and Memtable

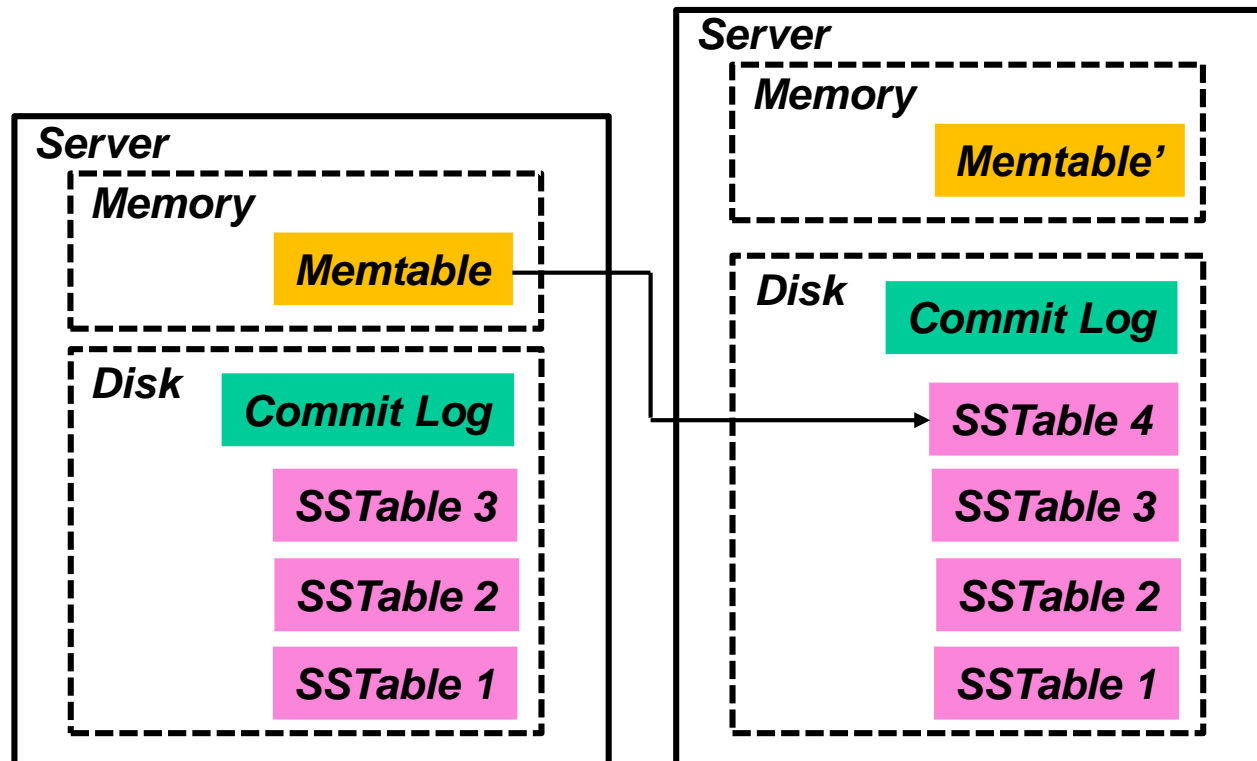
## Read Path



To optimize reads and to read relevant SSTables only, Bloom filters are used (★)

# LSM Trees - Flushing

When a Memtable reaches a certain size, it is frozen, a new Memtable is created, and the old Memtable is flushed in a new SSTable on disk



# The Write Path of an Insert

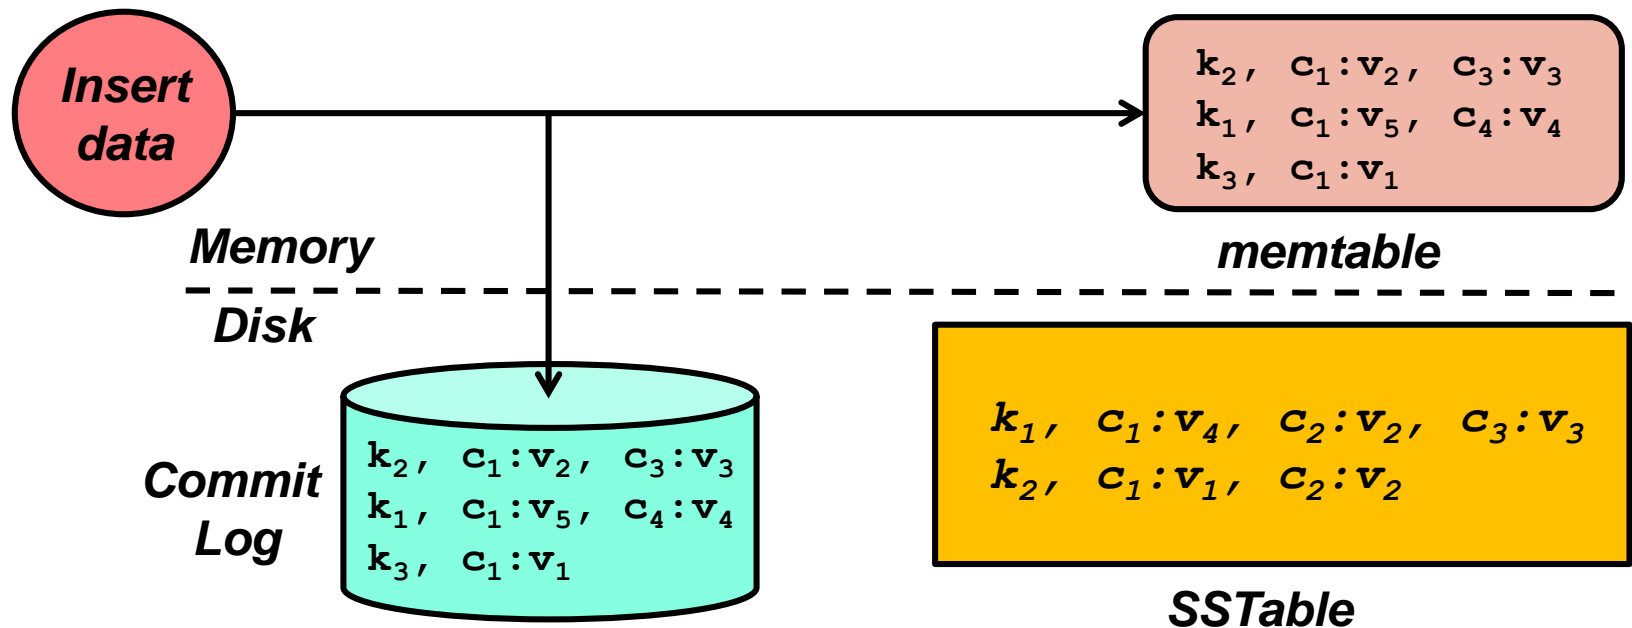
Assume row keys  $k_1$ ,  $k_2$ , and  $k_3$  map to the same partition

Assume rows  $(k_1, c_1:v_4, c_2:v_2, c_3:v_3)$  and

$(k_2, c_1:v_1, c_2:v_2)$  are already flushed in a SSTable on disk

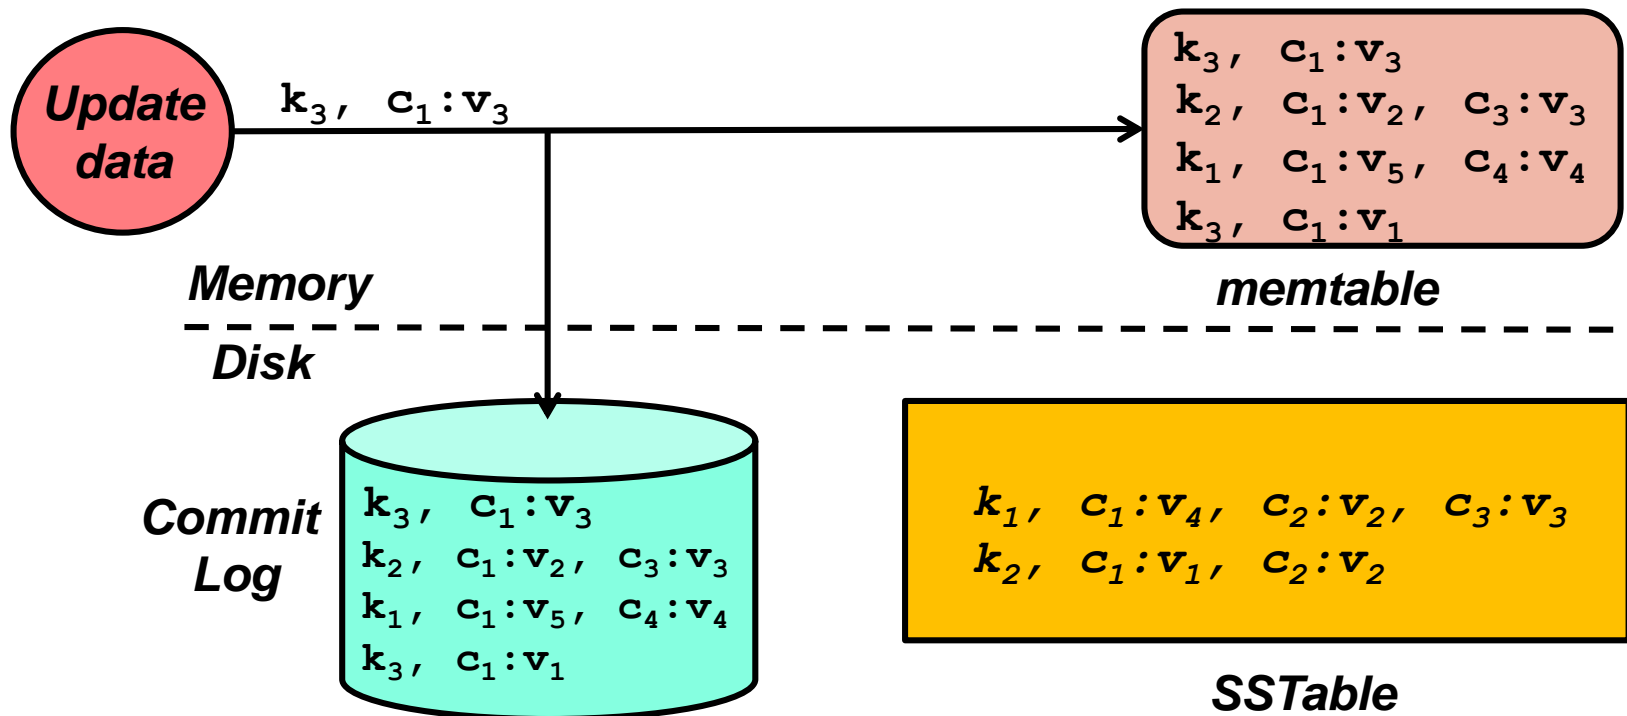
Next are rows  $(k_3, c_1:v_1)$ ,  $(k_1, c_1:v_5, c_4:v_4)$  and

$(k_2, c_1:v_2, c_3:v_3)$  written into Commit Log and memtable



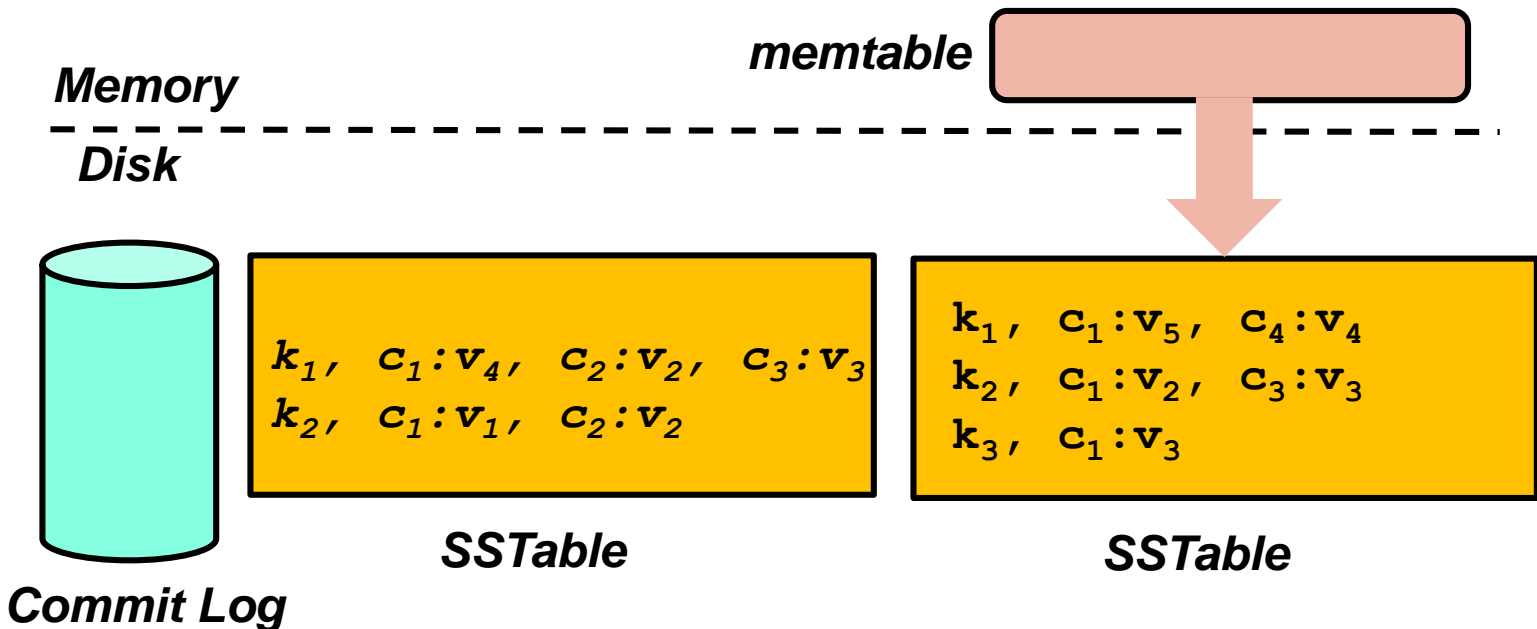
# The Write Path of an Update

- The update command works like an upsert:
  - It simply inserts a new row into the commit log and memtable,
  - Cassandra doesn't modify a column value in place



# Memtable Flushing

- When a memtable exceeds a configurable threshold, the memtable data are sorted by the primary key and flushed into a SSTable on disk
  - Only the latest value of each column of a row goes to the SSTable, since it has the greatest time stamp value
  - After flushing the memtable, the Commit Log data is also purged





# ***Log Structured Merge Trees (Summary)***

---

- Cassandra uses Log Structured Merge (LSM) Trees in a similar way as BigTable does:
  - Writes are done into a commit log on disk and in a memtable in memory,
  - Each column family (table) has its own memtable, commit log, and SSTables in each partition,
  - When the size of a memtable reaches a prescribed threshold, it is flushed in a SSTable on disk,
  - SSTables are immutable, hence different SSTables may contain different versions of a row column, and updates and deletes are implemented as time stamped writes (there is no in-place updates or deletes),
  - There are no reads before writes and no guarantee of the uniqueness of the primary key (unless special mechanisms are applied),
  - Reads are performed by merging requested data from the memtable and all SSTables,
    - To read only from SSTables containing data requested, Bloom Filters are used,
  - Clients are supplied the latest versions of data read,
  - SSTables are periodically compacted into a new SSTable

# ***Updates and Timestamps***

---

- Cassandra flushes only the most recent value of each column of a row in memtable to the SSTable
  - The most recent column value has the greatest timestamp
- Precise timestamps are needed if updates are frequent
- Timestamps are provided by clients
- Clocks of all client machines should be synchronized using NTP (network time protocol)

# About Reads

---

- Cassandra must combine results from the memtable and potentially multiple SSTables to satisfy a read
- First, Cassandra checks the Bloom filter
  - Each SSTable has a Bloom filter associated with it that checks the probability of having any data for the requested partition key in the SSTable before doing any disk I/O
- If the Bloom filter does not rule out the SSTable, Cassandra checks the partition key cache and takes a number of other actions to find the row fragment with the given key

# Bloom Filter

(1)

- A Bloom filter is a space-efficient probabilistic data structure that is used to test whether an element is a member of a set
  - False positive matches are possible, but
  - False negatives are not
- An empty Bloom filter is a  $m$  bit array, all bits set to 0
- There must also be  $n$  ( $< m$ ) different hash functions, each of which maps an element to one of the  $m$  array positions
- To add an element into the Bloom filter, its  $n$  array positions are calculated
  - Bits at all  $n$  positions are set to 1

# Bloom Filter

(2)

- To test whether an element is in the set, its  $n$  array positions are calculated
  - If any of the bits at these positions are 0, the element is definitely not in the set
  - Otherwise, it is probably in the set
- It has been proved that fewer than 10 bits per an element in the set are required for a 1% false positive probability, independent of the size or number of elements in the set

# Bloom Filter Example

- Assume:
  - The set of SSTable keys is {173, 256, 314}
  - Hash functions are:  $h_1 = k \bmod 7$ ,  $h_2 = k \bmod 13$ ,  $h_3 = k \bmod 17$ , where  $k$  is a SSTable key, and  $n = 3$
  - Let  $m = 17$
- Array positions:
  - For  $k = 173$ , array positions are (5, 4, 3)
  - For  $k = 256$ , array positions are (4, 9, 1)
  - For  $k = 314$ , array positions are (6, 2, 8)

0	1	1	1	1	1	1	0	1	1	0	0	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16

- Tests:
  - For  $k = 105$ , array positions are (0, 1, 3), so negative
  - For  $k = 106$ , array positions are (1, 2, 4), so false positive

# About Deletes

---

- CQL `DELETE` statement works as a write
  - Deleted column value is written in the memtable and commit log as a (row\_key, column\_name, TOMBSTONE, time\_stamp) tuple
    - TOMBSTONE is the new column value, indicating it has been deleted
- Data in a Cassandra column can have an optional expiration time called TTL (time to live)
- The TTL in seconds is set in CQL
  - Cassandra marks TTL data with a tombstone after the requested amount of time has expired
  - Tombstones exist for a period of time defined by `gc_grace_seconds` that is a table property (10 days by default)
  - After data is marked with a tombstone and `gc_grace_seconds` has elapsed, the data is automatically removed during the normal compaction process

# Storage Engine Example

**(1)**

row key user_id	column	SSTable 1	SSTable 2	SSTable 3	merge
jbond	name	James ts 10			James ts 10
	city	London ts 10		Paris ts 30	Paris ts 30
	email		jbond@ecs ts 20		jbond@ecs ts 20
	pet		dog ts 20	tombstone ts 40	tombstone ts 40



# Storage Engine Example (2)

---

- Assume:

```
(row_key: jbond, (name, James, ts10), (city, London, ts10))
```

has been flushed into SSTable 1

- After issuing CQL commands:

```
ALTER TABLE users ADD email text, pet text;
```

```
UPDATE users SET email='jbond@ecs.vuw.ac.nz',  
pet = 'dog' WHERE id = 'jbond';
```

the record

```
(row_key: jbond, (email, jbond@ecs.vuw.ac.nz,  
ts20), (pet, dog, ts20))
```

has been later flushed into SSTable 2

# Storage Engine Example

(3)

- Assume CQL commands:

```
UPDATE users SET city = 'Paris'
WHERE id = 'jbond';
```

```
DELETE pet FROM users WHERE id = 'jbond';
```

induce storing the following records into SSTable 3

```
(row_key: jbond, (city, Paris, ts30))
(row_key: jbond, (pet, tombstone, ts40))
```

- The command

```
READ * FROM users WHERE user_name = 'jbond';
```

returns:

id	name	city	email
jbond	James	Paris	jbond@ecs.vuw.ac.nz

# Compaction Strategies

---

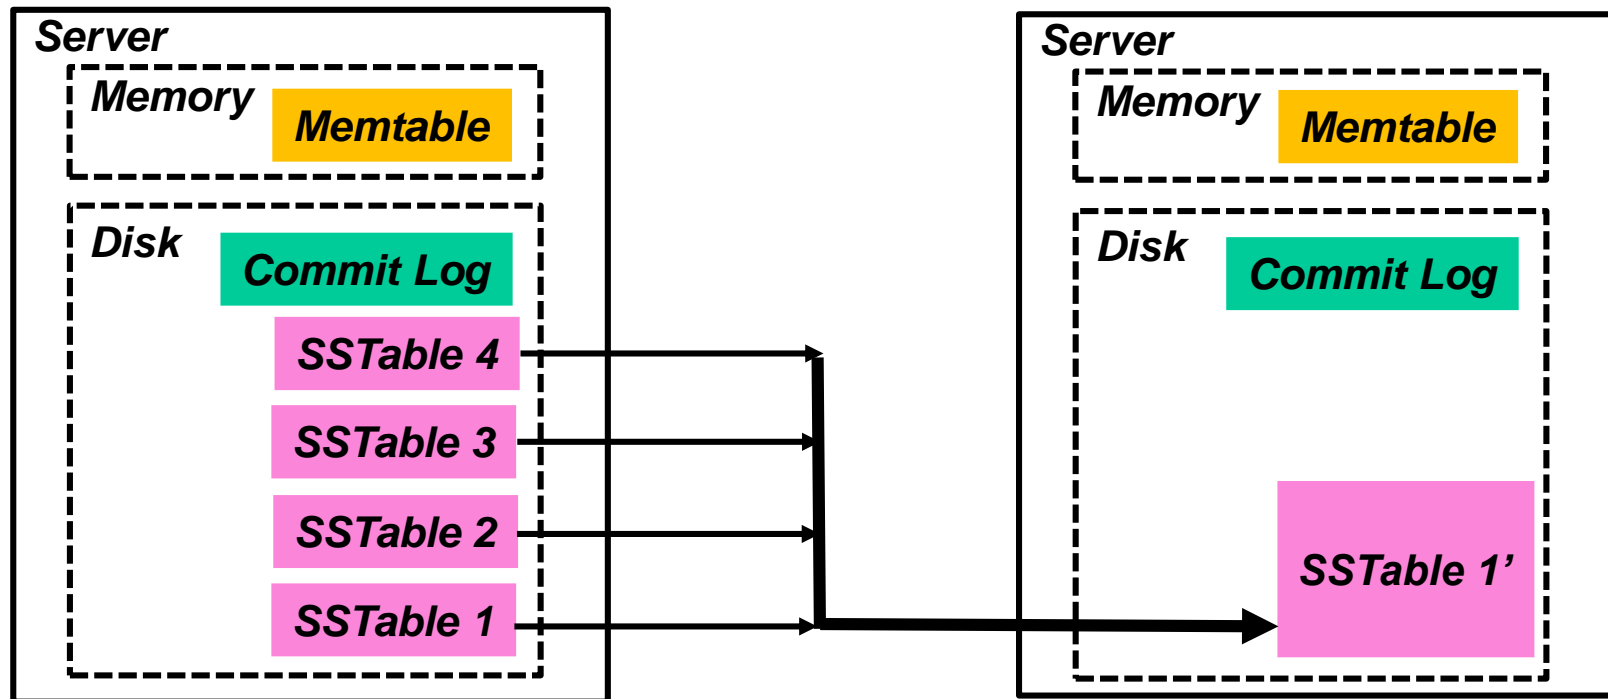
- Compaction is instrumental for attaining fast reads
- Cassandra supports:
  - `SizeTieredCompactionStrategy` designed for write intensive workloads (**default**),
  - `DateTieredCompactionStrategy` designed for time-series and expiring data, and
  - `LeveledCompactionStrategy` designed for read intensive workloads
- Compaction strategy is defined per column family and applied to its SSTables
- Cassandra performs compaction automatically, or it can be started manually using the `nodetool compact` command

# LSM Trees - Compaction

- Since SSTables are immutable, column value updates and deletes are accomplished solely by writes in the memtable:

**write**(old\_row\_key, old\_column\_key, new\_column\_value, new\_timestamp)

- In the case of deletes, the *new\_column\_value* is called the **tombstone**.
- To reclaim the disk space and speed up reads, after a certain number of memtables has been flushed, the compaction of SSTables is undertaken



# Size Tiered Compaction

**STCS produces SSTables of increasing size on different levels**

- The size of a SSTable is  $4^i m$ , where  $m$  is the size of the memtable and  $i$  is the level number ( $i = 0, 1, \dots$ )

**Size of a SSTable:**




- Problems in an update intensive workload:
  1. Read performance can be inconsistent since a row may have columns in many SSTables,
  2. Removing obsolete columns (particularly deleted ones) is a problem
  3. Time and storage space to do the compaction of bigger SSTables rise

# Leveled Compaction Strategy (LCS)

- LCS creates SSTables of a fixed size and groups them into levels  $L_1, L_2, L_3, \dots$ 
  - Each level  $L_i$  ( $i > 1$ ) has ten times greater size than the previous one and contains accordingly up to 10 times more SSTables
- The LCS compaction produces SSTables of a level by making a long sequential file that is split in a number of fixed size SSTables
- Accordingly, each row key appears on a single level at most once
  - There may exist an overlap of row keys on different levels
- Even more, with LCS:
  - The probability that only one of all SSTables contains a given row key value is  $\sim 0.9$ , and
  - The expected number of SSTables to contain a given row key value is  $\sim 1.1$

# LCS versus STCS – SSTables to Read



# WHEN to Use LCS

---

- LCS is a better option when:
  - Low latency reads are needed,
  - High read/write ratio,
  - Update intensive workload

- To declare LCS for a table:

```
CREATE TABLE <table_name> (...)  
WITH COMPACTION = { 'class' :  
  'LeveledCompactionStrategy' };
```

- In the `blogs` keyspace:

```
sqlsh.blogs => CREATE TABLE user_subs  
(user_name text PRIMARY KEY, no_of_subs  
counter) WITH COMPACTION = { 'class' :  
  'LeveledCompactionStrategy' };
```



# ***WHEN to Use STCS***

---

- STCS is default for a column family
- STCS is a better option than LCS if:
  - DISK I/O is expensive,
  - The workload is write heavy, and
  - Rows are written once (and rarely updated)
- In the `blogs` keyspace example, all tables except the `user_subs` and `suscribes_to` tables should have **Seize Tiered Compaction Strategy**
- Even the `blog_entries` table should be compacted by STCS since it is insert and not update intensive
  - Inserts write new rows
  - Updates write existing rows with a new column value

# Summary

---

- Cassandra's Storage Engine uses Log Structured Merge (LSM) Trees with:
  - A commit log on disk,
  - A per column family memtable in memory
- Memtables are flushed in immutable SSTables on disk
- Updates and deletes are implemented as time stamped writes
- Different SSTables may contain different versions of a row column
- SSTables are periodically compacted into a new SSTable